Qwen2.5VL

> **Demo Environment**

Development Board: RDK S100(P) Motherboard

Qwen2.5-VL is Qwen's new flagship visual language model, representing a significant leap forward compared to the previous Qwen2-VL.

1. Model Size

Model	Size
qwen2.5vl:3b	3.2GB
qwen2.5vl:7b	6.0GB

The 7B model requires 8GB of available runtime memory, while the RDK S100(P) motherboard only has a maximum available memory of about 6.7GB. Therefore, only the 3B model was used here.

2. Performance

		Qwen2.5-VL-3B	InternVL2.5-4B	Qwen2-VL-7B
College-level Problems	MMMU	53.1	52.3	**54.1**
	MMMU Pro	**31.6**	32.7	30.5
Document and Diagrams Reading	DocVQA	93.9	91.6	**94.5**
	InfoVQA	**77.1**	72.1	76.5
	CC-OCR	**74.5**	-	61.6
	OCRBenchV2	**54.3**	-	47.8
General Visual Question Answering	MegaBench	28.9	-	**34.4**
	MMStar	55.8	58.3	**60.7**
	MMBench1.1	77.6	79.3	**80.7**
Math	MathVista	**62.3**	60.5	58.2
	MathVision	**21.2**	20.9	16.3
Video Understanding	VideoMME	61.5	62.3	**63.3**
	MMBench-Video	1.6	**1.7**	1.4
	LVBench	**43.3**	-	-
	CharadesSTA	**38.8**	-	-
Visual Agent	AITZ	**76.9**	-	-
	Android Control	**63.7**	-	-
	ScreenSpot	**55.5**	-	55.3
	ScreenSpot Pro	**23.9**	-	1.6

		Qwen2.5-VL 7B	Qwen2-VL 7B	GPT-4o Mini	Other Best Open LVLM With Similar Size
College-level Problems	MMMU	58.6	54.1	**60.0**	56.0
	MMMU Pro	**38.3**	30.5	37.6	34.3
Document and Diagrams Reading	DocVQA	**95.7**	94.5	-	93.0
	InfoVQA	**82.6**	76.5	-	77.6
	CC-OCR	**77.8**	61.6	-	61.6
	OCRBenchV2	**56.3**	47.8	43	47.8
General Visual Question Answering	MegaBench	36.8	34.4	**43.1**	34.4
	MMStar	**63.9**	60.7	54.8	62.8
	MMBench1.1	**82.7**	80.7	76.0	79.4
Math	MathVista	**68.2**	58.2	52.5	67.2
	MathVision	**25.1**	16.3	-	19.7
Video Understanding	VideoMME	**65.1**	63.3	64.8	63.3
	MMBench-Video	**1.8**	1.4	-	1.7
	LVBench	**45.3**	-	-	38.4
	CharadesSTA	43.6	-	-	**48.4**
Visual Agent	AITZ	**81.9**	-	-	53.3
	Android Control	**60.1**	-	-	**61.5**
	ScreenSpot	84.7	55.3	-	**89.5**
	ScreenSpot Pro	29.0	1.6	-	**35.7**

3. Using Qwen2.5VL

3.1 Running Qwen2.5VL

Use the `run` command to start running the model. If you haven't downloaded this model before, it will automatically pull the model from the Ollam model library:

```
ollama run qwen2.5vl:3b
```

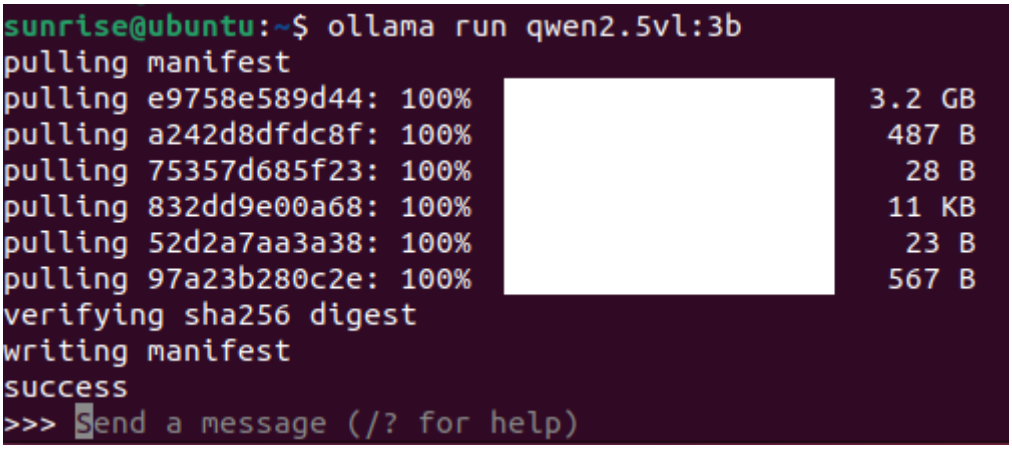

3.2 Engaging in a Dialogue

```
Please tell me how many hours there are in a day
```

The response time depends on your hardware configuration. Please be patient!

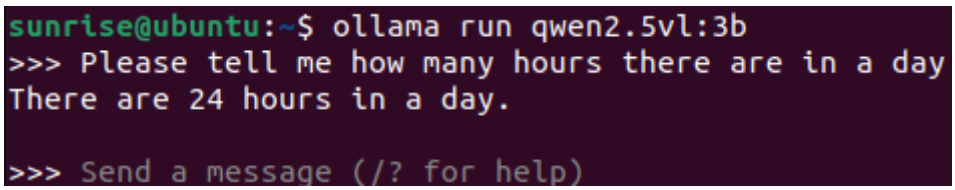

3.3 Visual Functions

```
What do you see in this picture? :./test_pic.png
#Using ": + image path" in the dialogue allows the model to use its visual
functions to interpret information from the image.
```

```
sunrise@ubuntu:~$ ollama run qwen2.5vl:3b
>>> What do you see in this picture?:./test_pic.png
Added image './test_pic.png'
The picture shows a person sitting on a rooftop or elevated platform,
looking up towards the sky. The person is wearing a white tank top and
appears to be in a relaxed or contemplative pose. In the foreground, there
is a clear plastic bottle. The background features a bright, sunny sky
with some clouds and a cityscape with buildings visible in the distance.
The overall atmosphere of the image is calm and serene.
```

3.4 Ending the Dialogue

Use the shortcut `Ctrl+d` or `/bye` to end the dialogue!

3.5 Chinese Dialogue

For those without a Chinese input method, please refer to the Chinese input method switching tutorial.

Chinese Dialogue:

```
>>> 一个小时是多少分钟？请用中文回答我
一个小时是60分钟。

>>> Send a message (/? for help)
```

References

> **Ollama**

Official Website: https://ollama.com/

GitHub: https://github.com/ollama/ollama

> **Qwen2.5VL**

GitHub: https://github.com/QwenLM/Qwen2.5-VL

Ollama Corresponding Model: https://ollama.com/library/qwen2.5vl