

Llava

Llava

1. Model Size
2. Performance
3. Using Llava
 - 3.1 Running Llava
 - 3.2 Initiating a Dialogue
 - 3.3 Visual Functions
 - 3.4 Ending a Dialogue
 - 3.5 Chinese Dialogue
- References


Demo Environment

Development Board: RDK S100(P) Motherboard

1. Model Size

Model	Size
llava:7b	4.7GB

2. Performance



3. Using Llava

3.1 Running Llava

Use the run command to start running the model. If this model has not been downloaded, it will automatically pull the model from the Ollama model library:

```
ollama run llava:7b
```

```
sunrise@ubuntu:~$ ollama run llava:7b
pulling manifest
pulling 170370233dd5: 100% [REDACTED] 4.1 GB
pulling 72d6f08a42f6: 100% [REDACTED] 624 MB
pulling 43070e2d4e53: 100% [REDACTED] 11 KB
pulling c43332387573: 100% [REDACTED] 67 B
pulling ed11eda7790d: 100% [REDACTED] 30 B
pulling 7c658f9561e5: 100% [REDACTED] 564 B
verifying sha256 digest
writing manifest
success
>>> Send a message (/? for help)
```

3.2 Initiating a Dialogue

```
How many minutes is half an hour?
```

Response time depends on hardware configuration; please be patient!

```
>>> How many minutes is half an hour?
Half an hour is 30 minutes.

>>> Send a message (/? for help)
```

3.3 Visual Functions



```
What do you see in this picture? ./test_pic.png
# Using ": + image path" in the dialogue allows the model to use its visual
functions to interpret information from the image.
```

```
sunrise@ubuntu:~$ ollama run llava:7b
>>> What do you see in this picture? ./test_pic.png
Added image './test_pic.png'
In the image, there is a man sitting on what appears to be a rooftop under a blue sky with some clouds. He has his head turned away from the camera and seems to be looking at something or someone out of frame. The man is wearing a tank top and shorts, and he is not wearing a shirt. Behind him, there's another person standing up, also facing away from the camera. Both individuals are on a rooftop with an open sky, suggesting a warm, sunny day. There is no visible text or distinctive branding in this image that I can discern.

>>> Send a message (/? for help)
```

3.4 Ending a Dialogue

Use the shortcut `ctrl+d` or `/bye` to end the dialogue!

3.5 Chinese Dialogue

For those without a Chinese input method, please refer to the [Chinese input method switching tutorial](#).

Chinese Dialogue:

```
>>> 一个小时是多少分钟？请用中文回答我
一小时是60分钟。

>>> Send a message (/? for help)
```

References

Ollama

Official Website: <https://ollama.com/>

GitHub: <https://github.com/ollama/ollama>

Llava

Ollama's corresponding model: <https://ollama.com/library/llava>