

# Multimodal Autonomous Agent Applications

---

## Multimodal Autonomous Agent Applications

1. Concept Introduction
  - 1.1 What is an "Autonomous Agent"?
  - 1.2 Implementation Principle Overview
2. Code Analysis
  - Key Code
    1. Agent Core Workflow (`largetmodel/utils/ai_agent.py`)
    2. Task Planning and LLM Interaction (`largetmodel/utils/ai_agent.py`)
    3. 参数处理与数据流实现 (`largetmodel/utils/ai_agent.py`)
  - Code Analysis
3. Practical Operation
  - 3.1 Configuring the Offline Large Model
    - 3.1.1 Configuring the LLM Platform (`yahboom.yaml`)
    - 3.1.2 Configure the model interface (`large_model_interface.yaml`)
    - 3.1.3 Recompile
  - 3.2 Start and Test the Functionality (Text Input Mode)
4. Common Problems and Solutions
  - 4.1 Abnormal Agent Behavior

Problem 1: The agent gets stuck in an infinite loop or repeatedly executes the same tool.
  - 4.2 Tool Invocation Failure

Problem 2: The agent correctly plans the action, but the tool execution fails.

---

## 1. Concept Introduction

### 1.1 What is an "Autonomous Agent"?

In the `largetmodel` project, the **multimodal autonomous agent** is the highest level of intelligence. It no longer simply responds to a user's command once, but is able to **autonomously think, plan, and continuously invoke multiple tools to complete a task to achieve a complex goal**.

The core of this functionality is the `agent_call` tool or its underlying \*\*toolchain manager (`ToolChainManager`). The autonomous agent is activated when a user makes a complex request that cannot be completed with a single tool call.

### 1.2 Implementation Principle Overview

The autonomous agent implementation in `largetmodel` follows the industry-leading **ReAct (Reason + Act)** paradigm. Its core idea is to mimic the human problem-solving process, cycling between "thinking" and "acting".

1. **Reason:** When the agent receives a complex goal, it first invokes a powerful Language Model (LLM) to "think." It asks itself, "What should my first step be to achieve this goal? Which tool should I use?" The LLM's output is not a final answer, but an action plan.
2. **Act:** Based on the LLM's reasoning, the agent executes the corresponding action—calling `ToolsManager` to run the specified tool (such as `visual_positioning`).

3. **Observe:** The agent retrieves the result of the previous action ("observation"), for example, `{"result": "The cup was found, located at [120, 300, 180, 360]"}.`
4. **Rethink:** The agent submits the observation results, along with the original goal, back to the LLM for a second round of "reasoning." It asks itself, "I've found the cup's location. What should I do next to find out its color?" The LLM might generate a new action plan, such as `{"thought": "I need to analyze the image of the area where the cup is located to determine its color", "action": "seewhat", "args": {"crop_area": [120, 300, 180, 360]}},`.

This **think -> act -> observe** loop continues until the initial goal is achieved, at which point the agent generates and outputs the final answer.

## 2. Code Analysis

---

### Key Code

#### 1. Agent Core Workflow (`largemodel/utils/ai_agent.py`)

The `_execute_agent_workflow` function is the Agent's main execution loop, defining the core "planning -> execution" process.

```
# From largemodel/utils/ai_agent.py

class AIAgent:
    # ...

    def _execute_agent_workflow(self, task_description: str) -> Dict[str, Any]:
        """
        Executes the agent workflow: Plan -> Execute.
        """
        try:
            # Step 1: Task Planning
            self.node.get_logger().info("AI Agent starting task planning phase")
            plan_result = self._plan_task(task_description)

            # ... (If the plan fails, return early)

            self.task_steps = plan_result["steps"]

            # Step 2: Perform all steps in sequence.
            execution_results = []
            tool_outputs = []

            for i, step in enumerate(self.task_steps):
                # 2.1. Before execution, process the data references in the
                # parameters.
                processed_parameters =
                self._process_step_parameters(step.get("parameters", {}), tool_outputs)
                step["parameters"] = processed_parameters

                # 2.2. Perform a single step
                step_result = self._execute_step(step, tool_outputs)
                execution_results.append(step_result)

                # 2.3. If the step is successful, save its output for reference
                # in subsequent steps.
        
```

```

        if step_result.get("success") and
step_result.get("tool_output"):
            tool_outputs.append(step_result["tool_output"])
        else:
            # If any step fails, abort the entire task.
            return { "success": False, "message": f"Task terminated
because step '{step['description']}' failed." }

        # ... Summarize and return the final results
        summary = self._summarize_execution(task_description,
execution_results)
        return { "success": True, "message": summary, "results": execution_results }

    # ... (Exception handling)

```

## 2. Task Planning and LLM Interaction (`largemodel/utils/ai_agent.py`)

The core of the `_plan_task` function is to construct a sophisticated Prompt, leveraging the large model's own inference capabilities to generate a structured execution plan.

```

# From largemodel/utils/ai_agent.py

class AIAgent:
    # ...
    def _plan_task(self, task_description: str) -> Dict[str, Any]:
        """
        Uses the large model for task planning and decomposition.
        """

        # Dynamically generate a list of available tools and their descriptions.
        tool_descriptions = []
        for name, adapter in
self.tools_manager.tool_chain_manager.tools.items():
            # ... (Retrieves tool description from adapter.input_schema)
            tool_descriptions.append(f"- {name}({{params}}): {{description}}")
        available_tools_str = "\n".join(tool_descriptions)

        # Building a highly structured planning Prompt
        planning_prompt = f"""
***# 可用工具:**
{available_tools_str}

***# 核心规则:**
1. **数据传递**: 当后续步骤需要使用之前步骤的输出时, **必须**使用
`{{{{steps.N.outputs.KEY}}}}` 格式进行引用。
    - `N` 是步骤的ID (从1开始)。
    - `KEY` 是之前步骤输出数据中的具体字段名。
2. **JSON格式**: 必须严格返回JSON对象。

***# 用户任务:**
{task_description}
"""

        # calling large models for planning

```

```

        messages_to_use = [{"role": "user", "content": planning_prompt}]
        # Note that this calls a generic text reasoning interface.
        result = self.node.model_client.infer_with_text("", 
message=messages_to_use)

        # ... (Parse the JSON response and return a list of steps)

```

### 3. 参数处理与数据流实现 (`largeModel/utils/ai_agent.py`)

`_process_step_parameters` 函数负责解析占位符，实现步骤间的数据流动。

```

# From largeModel/utils/ai_agent.py

class AIAgent:
    # ...
    def _process_step_parameters(self, parameters: Dict[str, Any],
previous_outputs: List[Any]) -> Dict[str, Any]:
        """
        Parses parameter dictionary, finds and replaces all {{...}} references.
        """

        processed_params = parameters.copy()
        # The regular expression is used to match placeholders in the format
        # {steps.N.outputs.KEY}.
        pattern = re.compile(r"\{\{(\w+)(\.\w+)*\.\w+\}\}")

        for key, value in processed_params.items():
            if isinstance(value, str) and pattern.search(value):
                # Use `re.sub` and a replacement function to process all found
                # placeholders
                # The replacement function searches the `previous_outputs` list
                # and returns the value.
                processed_params[key] = pattern.sub(replacer_function, value)

        return processed_params

```

## Code Analysis

The AI Agent is the "central brain" of the system. It transforms the high-level, and sometimes even fuzzy, tasks posed by the user into a series of precise and ordered functions. The Agent is a multi-step, scalable platform that does not rely on any specific model platform. Instead, it is built on a general, extensible architecture.

- **Dynamic Task Planning:** The core capability of the Agent lies in the `_plan_task` function. It does not rely on hard-coded logic but dynamically generates task plans through interaction with a large model.
- **Self-Awareness and Prompt Construction:** At the start of planning, the Agent first examines all available tools and their descriptions. Then, it packages this tool information, user tasks, and strict rules (such as data transmission formats) into a highly structured `planning_prompt`.
- **Model as Planner:** This Prompt is sent to a general text-based large model. The model infers based on the provided context and returns a multi-step action plan in JSON format. This design is highly scalable: when tools are added or modified in the system, the Agent's planning capabilities are automatically updated without requiring modifications to the Agent's own code.

**2. Toolchain and Data Flow:** Real-world tasks often require multiple tools to collaborate. For example, "taking a picture and describing it" requires the output (image path) of the "taking the picture" tool to be used as the input of the "describing" tool. The AI Agent elegantly achieves this through the `_process_step_parameters` function.

- **Data Reference Placeholders:** During the planning phase, the large model embeds special placeholders, such as `{{steps.1.outputs.data}}`, into the parameter values that need to transmit data.
- **Real-time Parameter Replacement:** In the main loop of `_execute_agent_workflow`, `_process_step_parameters` is called before each step. It uses regular expressions to scan all parameters of the current step. Once a placeholder is found, it retrieves the corresponding data from the output list of the previous step and replaces it in real time. This mechanism is key to automating complex tasks.

**3. Supervised Execution and Fault Tolerance:** `_execute_agent_workflow` constitutes the Agent's main execution loop. It strictly follows the planned sequence of steps, executing each action sequentially and ensuring correct data transfer between steps.

- **Atomic Steps:** Each step is treated as an independent "atomic operation." If any step fails, the entire task chain immediately aborts and reports an error. This ensures system stability and predictability, preventing continued execution in an erroneous state.

In summary, the general implementation of the AI Agent demonstrates an advanced software architecture: it doesn't directly solve the problem, but rather builds a framework for an external, general-purpose inference engine (large model) to solve it. Through the two core mechanisms of "dynamic programming" and "data flow management," the Agent can orchestrate a series of independent tools into complex workflows capable of performing advanced tasks.

## 3. Practical Operation

### 3.1 Configuring the Offline Large Model

#### 3.1.1 Configuring the LLM Platform (`yahboom.yaml`)

This file determines which large model platform the `model_service` node loads as its primary language model.

1. Open the file in the terminal:

```
vim ~/yahboom_ws/src/largemode1/config/yahboom.yaml
```

2. Modify/Confirm `llm_platform`:

```
model_service:                      #Model server node parameters
  ros__parameters:
    language: 'en'                  #Large Model Interface Language
    useolinetts: True               #This option is invalid in text mode and
can be ignored.

# Large Model Configuration
llm_platform: 'ollama'             # Key point: Make sure this is 'ollama'
```

### 3.1.2 Configure the model interface (`large_model_interface.yaml`)

This file defines which visual model is used when the platform is selected as `ollama`.

1. Open the file in the terminal

```
vim ~/yahboom_ws/src/largemode1/config/large_model_interface.yaml
```

2. Locate the Ollama-related configuration

```
#.....  
## Offline Large Language Models  
# ollama Configuration  
ollama_host: "http://localhost:11434" # ollama server address  
ollama_model: "llava" # Key: Replace this with your downloaded multimodal model,  
such as "llava"  
#.....
```

**Note:** Ensure that the model specified in the configuration parameters (e.g., `llava`) can handle multimodal input.

### 3.1.3 Recompile

```
cd ~/yahboom_ws/  
colcon build  
source install/setup.bash
```

## 3.2 Start and Test the Functionality (Text Input Mode)

1. Start the `largemode1` main program (text mode):

Open a terminal and run the following command:

```
ros2 launch largemode1 largemode1_control.launch.py text_chat_mode:=true
```

2. Send a text command:

Open another terminal and run the following command:

```
ros2 run text_chat text_chat
```

Then start typing the text: "Generate an image similar to the current scene based on the current environment."

3. Observations:

In the first terminal running the main program, you will see log output showing that the system received a text command, invoked the `aiagent` tool, and then provided a prompt to the LLM. The LLM will analyze the detailed steps of the tool invocation. For example, this question will invoke the `seewhat` tool to obtain the image, then provide the image to the LLM for parsing. The parsed text will then be provided to the LLM as content for generating a new image.

## 4. Common Problems and Solutions

## 4.1 Abnormal Agent Behavior

**Problem 1: The agent gets stuck in an infinite loop or repeatedly executes the same tool.**

**Solutions:**

1. **Enhance the Prompt:** Add stronger restrictions to the ReAct Prompt, such as: "Do not repeat the same operation," and "If two consecutive observations are the same, try a different tool or end the task."
2. **Add Iteration Limits:** Set a maximum number of loop iterations, such as `max_turns` in the code example, to prevent infinite loops.
3. **Replace with a More Powerful LLM:** The agent's logical capabilities largely depend on the "intelligence" of the backend LLM. A more powerful model can better understand tasks and make plans.

## 4.2 Tool Invocation Failure

**Problem 2: The agent correctly plans the action, but the tool execution fails.**

**Solution:**

1. **Inspect Individual Tools:** This is usually not a problem with the agent logic, but rather with the tool being called (such as `seewhat` or `visual_positioning`). Please follow the tutorial for the corresponding tool to troubleshoot the problem.
2. **Error Handling:** A robust `try...except` mechanism is needed in the `ToolChainManager` loop to catch exceptions during tool execution and feed the error information back to the LLM as an "observation," letting it know that the previous step failed and allowing it to replan.